

CS 221 Section 5: Games

Jerry Qu & Will Deaderick

Section Contents

- Minimax Search, w/ example
- Evaluation Functions
- Alpha-Beta Pruning, w/ example
- Alpha-Beta practice problem
- Game Theory practice problem
- MCTS & AlphaZero (time permitting)

Minimax

- Applicable to two-player zero-sum games (and other adversarial games)
- Characteristics:
 - Players take turns
 - Utility comes from end state, no intermediate utility
- Overall goal is to model outcomes when both players play optimally
- When state transitions are not deterministic, we use *expectiminimax*.
- Complexity: $O(d)$ space, $O(b^{2d})$ time, for depth d , branching factor b

Evaluation Functions

- Goal is to provide some estimate of $V_{\text{minimax}}(s)$ for state s .
- Oftentimes requires domain knowledge since we can't make additional moves.
- Because of large time complexity for searching games, evaluation functions are often necessary.
- Pitfalls: in some cases where utility can fluctuate (e.g. a piece capturing sequence in chess), evaluation functions may become inaccurate.

Alpha-Beta Pruning: Interpretation

- At each “**Max**” state, the “**Max**” player keeps track of **alpha**, a lower bound on her value for that state.

Alpha-Beta Pruning: Interpretation

- At each “**Max**” state, the “**Max**” player keeps track of **alpha**, a lower bound on her value for that state.
 - Values lower than **alpha** don’t interest “**Max**”

Alpha-Beta Pruning: Interpretation

- At each “**Max**” state, the “**Max**” player keeps track of **alpha**, a lower bound on her value for that state.
 - Values lower than **alpha** don’t interest “**Max**”
- At each “**Min**” state, the “**Min**” player keeps track of **beta**, an upper bound on her value for that state.

Alpha-Beta Pruning: Interpretation

- At each “**Max**” state, the “**Max**” player keeps track of **alpha**, a lower bound on her value for that state.
 - Values lower than **alpha** don’t interest “**Max**”
- At each “**Min**” state, the “**Min**” player keeps track of **beta**, an upper bound on her value for that state.
 - Values larger than **beta** don’t interest “**Min**”

Alpha-Beta Pruning: Rules

- Start with **alpha** as negative infinity, and **beta** as infinity
- Propagate **alpha** and **beta** down the search tree
- At each “**Max**” node, update **alpha** if we find a larger (than **alpha**) child leaf value, child alpha, or child beta
- At each “**Min**” node, update **beta** if we find a smaller (than **beta**) child leaf value, child alpha, and child beta
- If **alpha** \geq **beta**, prune!

Alpha-Beta Pruning: Intuition

- We can prune beneath a “**Max**” node if that “**Max**” node’s parent is already guaranteed a better (i.e. smaller) value than any of the “**Max**” node’s children.

Alpha-Beta Pruning: Intuition

- We can prune beneath a “**Max**” node if that “**Max**” node’s parent is already guaranteed a better (i.e. smaller) value than any of the “**Max**” node’s children.
- We can prune beneath a “**Min**” node if that “**Min**” node’s parent is already guaranteed a better (i.e. larger) value than any of the “**Min**” node’s children.

Extensions

- Beam search (“forward pruning”) - use an evaluation function to only expand the most promising child nodes, effectively lowering the branching factor

Extensions

- Beam search (“forward pruning”) - use an evaluation function to only expand the most promising child nodes, effectively lowering the branching factor
 - Pros: similar to human intelligence, computational efficiency

Extensions

- Beam search (“forward pruning”) - use an evaluation function to only expand the most promising child nodes, effectively lowering the branching factor
 - Pros: similar to human intelligence, computational efficiency
 - Cons: similar to human intelligence, no more guarantees

Extensions

- Beam search (“forward pruning”) - use an evaluation function to only expand the most promising child nodes, effectively lowering the branching factor
 - Pros: similar to human intelligence, computational efficiency
 - Cons: similar to human intelligence, no more guarantees
- Dynamic move ordering - Manually specify which moves to expand first, or use an evaluation function, or keep track of which moves have been most successful in the past

Game Theory

- Applicable to single-move simultaneous games (covered already) and more.
- Defined by a set of players, actions, and utility matrix.
- We can have *pure strategies* or *mixed strategies*.
- Nash Equilibrium: a (pure or mixed) strategy for all players such that no player can gain utility by changing their strategy.
 - In any finite-player game with finite actions, there's ALWAYS at least one Nash Equilibrium.